

Leaf Disease Classification - Keras CNN .. With 96% Accuracy

In [1]:

```
import warnings
warnings.filterwarnings("ignore")
import tensorflow as tf
import matplotlib.pyplot as plt
tf.compat.v1.set_random_seed(0)
from tensorflow import keras
import numpy as np
np.random.seed(0)
import itertools
from keras.preprocessing.image import image_dataset_from_directory
from tensorflow.keras.layers.experimental.preprocessing import Rescaling
from sklearn.metrics import precision_score, accuracy_score, recall_score, confusion_mat
```

Data Loading

Setting up Image Data Generators

In [2]:

```
#train_gen = image_dataset_from_directory(directory="../input/new-plant-diseases-dataset",
#test_gen = image_dataset_from_directory(directory="../input/new-plant-diseases-dataset/

train_gen = image_dataset_from_directory(directory="../input/new-plant-diseases-dataset/",
                                         image_size=(256, 256))
test_gen = image_dataset_from_directory(directory="../input/new-plant-diseases-dataset/N",
                                         image_size=(256, 256))

rescale = Rescaling(scale=1.0/255)
train_gen = train_gen.map(lambda image,label:(rescale(image),label))
test_gen = test_gen.map(lambda image,label:(rescale(image),label))
```

Found 70295 files belonging to 38 classes.

```
2022-06-02 09:49:51.830329: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:51.990947: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:51.991830: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:51.994026: I tensorflow/core/platform/cpu_feature_guard.cc:142] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2022-06-02 09:49:51.994343: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:51.995141: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:51.995803: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:54.246926: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:54.247814: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:54.248494: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-06-02 09:49:54.249073: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 15403 MB memory: -> device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute capability: 6.0
```

Found 17572 files belonging to 38 classes.

Data Engineering

Since the data is already augmented, there is no requirement of data engineering. Feature scaling is automatically done by image generators

Modelling

In [3]:

```
model = keras.Sequential()

model.add(keras.layers.Conv2D(32,(3,3),activation="relu",padding="same",input_shape=(256
model.add(keras.layers.Conv2D(32,(3,3),activation="relu",padding="same"))
model.add(keras.layers.MaxPooling2D(3,3))

model.add(keras.layers.Conv2D(64,(3,3),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(64,(3,3),activation="relu",padding="same"))
model.add(keras.layers.MaxPooling2D(3,3))

model.add(keras.layers.Conv2D(128,(3,3),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(128,(3,3),activation="relu",padding="same"))
model.add(keras.layers.MaxPooling2D(3,3))

model.add(keras.layers.Conv2D(256,(3,3),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(256,(3,3),activation="relu",padding="same"))

model.add(keras.layers.Conv2D(512,(5,5),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(512,(5,5),activation="relu",padding="same"))

model.add(keras.layers.Flatten())

model.add(keras.layers.Dense(1568,activation="relu"))
model.add(keras.layers.Dropout(0.5))

model.add(keras.layers.Dense(38,activation="softmax"))

opt = keras.optimizers.Adam(learning_rate=0.0001)
model.compile(optimizer=opt,loss="sparse_categorical_crossentropy",metrics=['accuracy'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 256, 256, 32)	896
conv2d_1 (Conv2D)	(None, 256, 256, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 85, 85, 32)	0
conv2d_2 (Conv2D)	(None, 85, 85, 64)	18496
conv2d_3 (Conv2D)	(None, 85, 85, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 28, 28, 64)	0
conv2d_4 (Conv2D)	(None, 28, 28, 128)	73856
conv2d_5 (Conv2D)	(None, 28, 28, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 9, 9, 128)	0
conv2d_6 (Conv2D)	(None, 9, 9, 256)	295168
conv2d_7 (Conv2D)	(None, 9, 9, 256)	590080
conv2d_8 (Conv2D)	(None, 9, 9, 512)	3277312
conv2d_9 (Conv2D)	(None, 9, 9, 512)	6554112
flatten (Flatten)	(None, 41472)	0
dense (Dense)	(None, 1568)	65029664
dropout (Dropout)	(None, 1568)	0
dense_1 (Dense)	(None, 38)	59622
=====		
Total params: 76,092,966		
Trainable params: 76,092,966		
Non-trainable params: 0		
=====		

In [4]:

```
ep = 10
history = model.fit_generator(train_gen,
                             validation_data=test_gen,
                             epochs = ep)
```

Epoch 1/10

2022-06-02 09:49:57.725110: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)

2022-06-02 09:50:00.285178: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version 8005

2197/2197 [=====] - 316s 140ms/step - loss: 1.6022 - accuracy: 0.5339 - val_loss: 0.5197 - val_accuracy: 0.8419

Epoch 2/10

2197/2197 [=====] - 195s 89ms/step - loss: 0.4302 - accuracy: 0.8633 - val_loss: 0.2852 - val_accuracy: 0.9066

Epoch 3/10

2197/2197 [=====] - 194s 88ms/step - loss: 0.2553 - accuracy: 0.9166 - val_loss: 0.2539 - val_accuracy: 0.9214

Epoch 4/10

2197/2197 [=====] - 195s 88ms/step - loss: 0.1754 - accuracy: 0.9418 - val_loss: 0.2177 - val_accuracy: 0.9314

Epoch 5/10

2197/2197 [=====] - 193s 88ms/step - loss: 0.1299 - accuracy: 0.9573 - val_loss: 0.1694 - val_accuracy: 0.9451

Epoch 6/10

2197/2197 [=====] - 194s 88ms/step - loss: 0.1014 - accuracy: 0.9661 - val_loss: 0.1315 - val_accuracy: 0.9583

Epoch 7/10

2197/2197 [=====] - 194s 88ms/step - loss: 0.0823 - accuracy: 0.9719 - val_loss: 0.1392 - val_accuracy: 0.9585

Epoch 8/10

2197/2197 [=====] - 204s 93ms/step - loss: 0.0691 - accuracy: 0.9769 - val_loss: 0.1477 - val_accuracy: 0.9547

Epoch 9/10

2197/2197 [=====] - 194s 88ms/step - loss: 0.0623 - accuracy: 0.9796 - val_loss: 0.1131 - val_accuracy: 0.9668

Epoch 10/10


2197/2197 [=====] - 194s 88ms/step - loss: 0.0532 - accuracy: 0.9829 - val_loss: 0.1166 - val_accuracy: 0.9677

Metrics

In [5]:

```
plt.figure(figsize = (20,5))
plt.subplot(1,2,1)
plt.title("Train and Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.plot(history.history['loss'],label="Train Loss")
plt.plot(history.history['val_loss'], label="Validation Loss")
plt.xlim(0, 10)
plt.ylim(0.0,1.0)
plt.legend()

plt.subplot(1,2,2)
plt.title("Train and Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.plot(history.history['accuracy'], label="Train Accuracy")
plt.plot(history.history['val_accuracy'], label="Validation Accuracy")
plt.xlim(0, 9.25)
plt.ylim(0.75,1.0)
plt.legend()
plt.tight_layout()
```



In [6]:

```
labels = []
predictions = []
for x,y in test_gen:
    labels.append(list(y.numpy()))
    predictions.append(tf.argmax(model.predict(x),1).numpy())
```

In [7]:

```
predictions = list(itertools.chain.from_iterable(predictions))
labels = list(itertools.chain.from_iterable(labels))
```

In [8]:

```
print("Train Accuracy : {:.2f} %".format(history.history['accuracy'][-1]*100))
print("Test Accuracy : {:.2f} %".format(accuracy_score(labels, predictions) * 100))
print("Precision Score : {:.2f} %".format(precision_score(labels, predictions, average='micro'))
print("Recall Score : {:.2f} %".format(recall_score(labels, predictions, average='micro'))
```

```
Train Accuracy : 98.29 %
Test Accuracy : 96.77 %
Precision Score : 96.77 %
Recall Score : 96.77 %
```


Confusion Matrix

In [9]:

```
plt.figure(figsize= (20,5))
cm = confusion_matrix(labels, predictions)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=list(range(1,39)))

fig, ax = plt.subplots(figsize=(15,15))
disp.plot(ax=ax,colorbar= False,cmap = 'YlGnBu')
plt.title("Confusion Matrix")
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

<Figure size 1440x360 with 0 Axes>



In []:

